

Independent Component Analysis (ICA)

1 Overview

The purpose of this assignment is to build an algorithm to successfully separate mixed signals (blind source separation). In other words, if we were to mix some source signals $\mathbf{U}(\mathbf{n}, \mathbf{t})$, where \mathbf{n} is the number of signals and \mathbf{t} is the length of each signal, with a mixing matrix $\mathbf{A}(\mathbf{m}, \mathbf{n})$, where \mathbf{m} is the number of mixed signals, we would be left with a mixed signals $\mathbf{X}(\mathbf{m}, \mathbf{t}) = \mathbf{A}\mathbf{U}$. Our goal is to recover the source signals (\mathbf{U}) with minimal information about the source data (\mathbf{U}). More formally, we are looking for an unmixing matrix $\mathbf{W}(\mathbf{n}, \mathbf{m}) = \mathbf{A}^{-1}$, such that $\mathbf{Y} = \mathbf{W}\mathbf{X}$ is the recovered signal. Independent component analysis, or ICA, is one such method to recover data with minimal or no information about the source data.

2 Method

First we can characterize the distributions of each source signal. Suppose each source signal U_j can be described by a density p_{U_j} ; the joint distribution of the source \mathbf{U} is given by:

$$p(\mathbf{U}) = \prod_{j=1}^n p_{U_j}(U_j).$$

Since the mixed signal \mathbf{X} is simply a linear transformation of the original signal \mathbf{U} , we can expect the distribution of the mixed signal \mathbf{X} to be similar to the distribution of the original signal \mathbf{U} . The distribution of the mixed signal \mathbf{X} is given by:

$$p(\mathbf{X}) = \prod_{j=1}^n p_{U_j}(\mathbf{W}\mathbf{X}) \cdot |\mathbf{W}|,$$

where $\mathbf{W} = \mathbf{A}^{-1}$.

As we can see from the distribution above, we have two unknowns: \mathbf{W} and p_{U_j} . We can *guess* as to what p_{U_j} may be by first defining a cumulative distribution function (cdf) for the original signal \mathbf{U} . Then, the distribution p_{U_j} is simply the derivative of that cdf. According to material by Andrew Ng [1], the sigmoid function $g(X) = 1/(1 + \exp(-X))$ is a reasonable *default* cdf; thus, $p_{U_j}(X) = g'(X)$. Any distribution can be used as long as it is non-Gaussian. Now, the only unknown is \mathbf{W} .

To solve for \mathbf{W} , we can use gradient ascent to maximize the log likelihood of \mathbf{W} (substitute $g'(X)$ for p_{U_j} and take the log of the distribution above):

$$l(\mathbf{W}) = \sum \left(\sum \log g'(\mathbf{W}\mathbf{X}) + \log |\mathbf{W}| \right).$$

Then, the gradient ascent algorithm can be described as:

$$W := W + \eta(I * t + (1 - 2 * Z) * Y^T) * W,$$

where η is the learning rate (1e-6), t is the length of each source signal, $Y = WX$ is the *recovered* data, and $Z = g(Y)$. After a set number of iterations ($R_max=1000$), we will have an approximate unmixing matrix W , such that $Y = WX$.

3 Experiments

The three variables I will be considering in this homework are: η (learning rate), R_max (number of iterations to update W), and m (the number of mixed signals). I will use the `numpy.linalg.norm` function and the Pearson correlation coefficient (`correlation_coefficient`) to determine how close the recovered signal Y is from the original signal U .

From initial trial-and-error, it was found large learning rates ($>1e-6$) would lead to an overflow at any iteration above 100; so, I tested learning rates $\leq 1e-6$. Moreover, I noticed the ICA algorithm struggled when it had to recover a larger number of signals. **Table 1** below indicates the correlation coefficient found for n recovered signals (`sounds.mat`), where n is between 2 and 5, $R_max=1000$, and $m=n$.

Table 1

n	Correlation Coefficient	Best η
2	0.9764	1e-6
3	0.7349	1e-7
4	0.5325	1e-8
5	0.3631	1e-9

The correlation coefficients seem to decrease with more signals. Moreover, smaller learning rates were required to obtain (relatively) better correlation coefficients. Since a learning rate of 1e-6 recovered approximately 97% of the first two signals, I will be using a learning rate of 1e-6 for future experiments.

Next, I found that the number of iterations (1000, 10000, 100000) produced very similar results and errors (`numpy.linalg.norm` and `correlation_coefficient`). The only difference between iterations was the run time to loop through each value of R_max ; larger values of R_max took longer time – 5sec, 60sec, and 300sec respectively. So, the learning rate was set to 1e-6 and R_max was set to 1000.


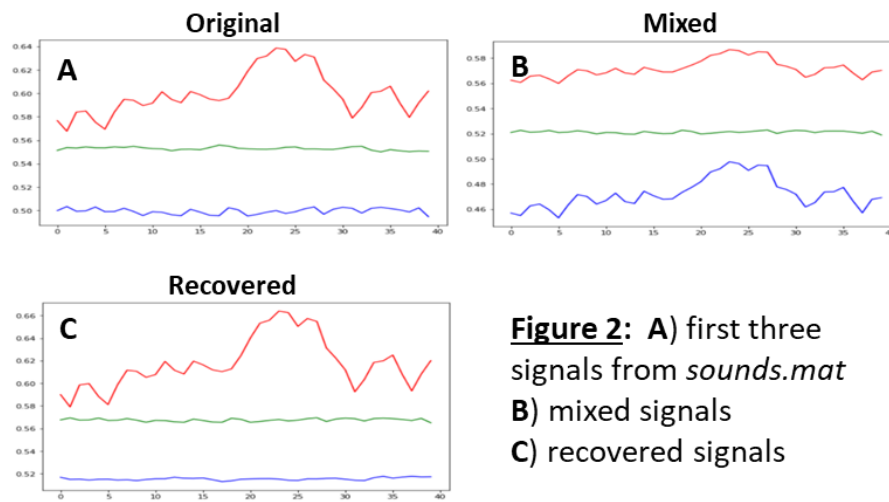


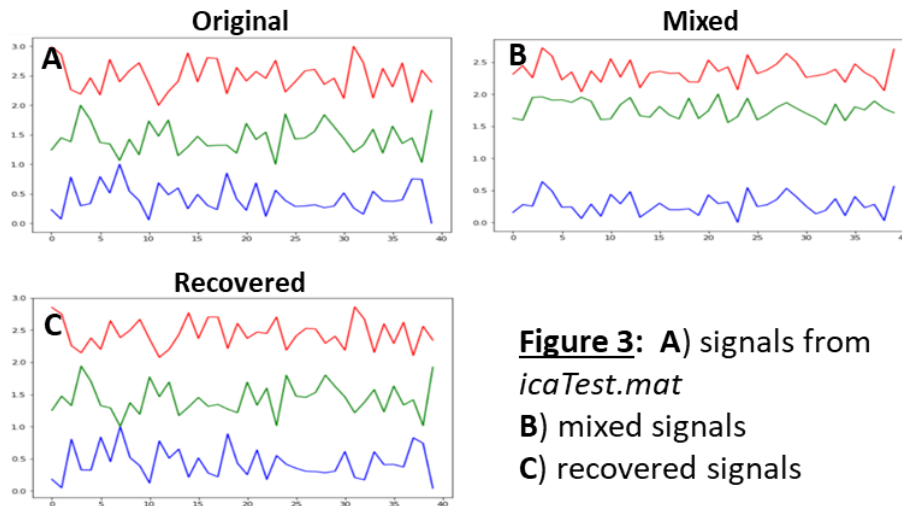
Figure 1: Norm error (`numpy.linalg.norm`) for values of m ranging from 5 to 20

Lastly, I experimented with the number of mixed signals m produced by the mixing matrix. **Figure 1** illustrates the error between the recovered and original signals. As seen by **Figure 1**, there is no clear pattern as to what values of m are optimal to reduce the recovered signal error. The lowest error was found when $m = n = 5$ (`sounds.mat`), and so for future experiments, m will be set to n .

After setting the learning rate to $1e-6$, R_max to 1000, and m to n , I mixed and recovered the signals provided by the *sounds.mat* file. **Figure 2** illustrates the original, mixed, and recovered signals.



As seen by **Figure 2**, the original and recovered signals look comparable. However, there were more promising results for the data provided by the *icaTest.mat* file. **Figure 3** illustrates the original, mixed, and recovered signals.



The original and recovered signal look very similar, and this is supported by the high correlation coefficient (**0.9796**).

4 Results and Discussion

Overall, the ICA algorithm was able to successfully recover approximately 98% of the ICA test data, with a learning rate of $1e-6$, an R_max of 1000, and with $m=n$. The program was prone to overflowing, thus requiring a very small learning rate. The number of iterations did not seem to improve the quality of the recovered images, indicating \mathbf{W} must be converging at some $R_max < 1000$. I observed the recovered signals would not always be in the same order as the original data. With further investigation, I found after rearranging the recovered signals to match the order of the original signals, the correlation coefficient would increase significantly; however, the `numpy.linalg.norm` function would be consistent with its results. The ICA algorithm was an interesting algorithm to learn and seemed to recover original data with limited information.

5 References

- a) `time` library – `time`
- b) `scipy` library – `io.loadmat`
- c) `numpy` library – `linalg.norm`
- d) `matplotlib` library – `pyplot`

[1] <http://cs229.stanford.edu/notes2020spring/cs229-notes11.pdf>